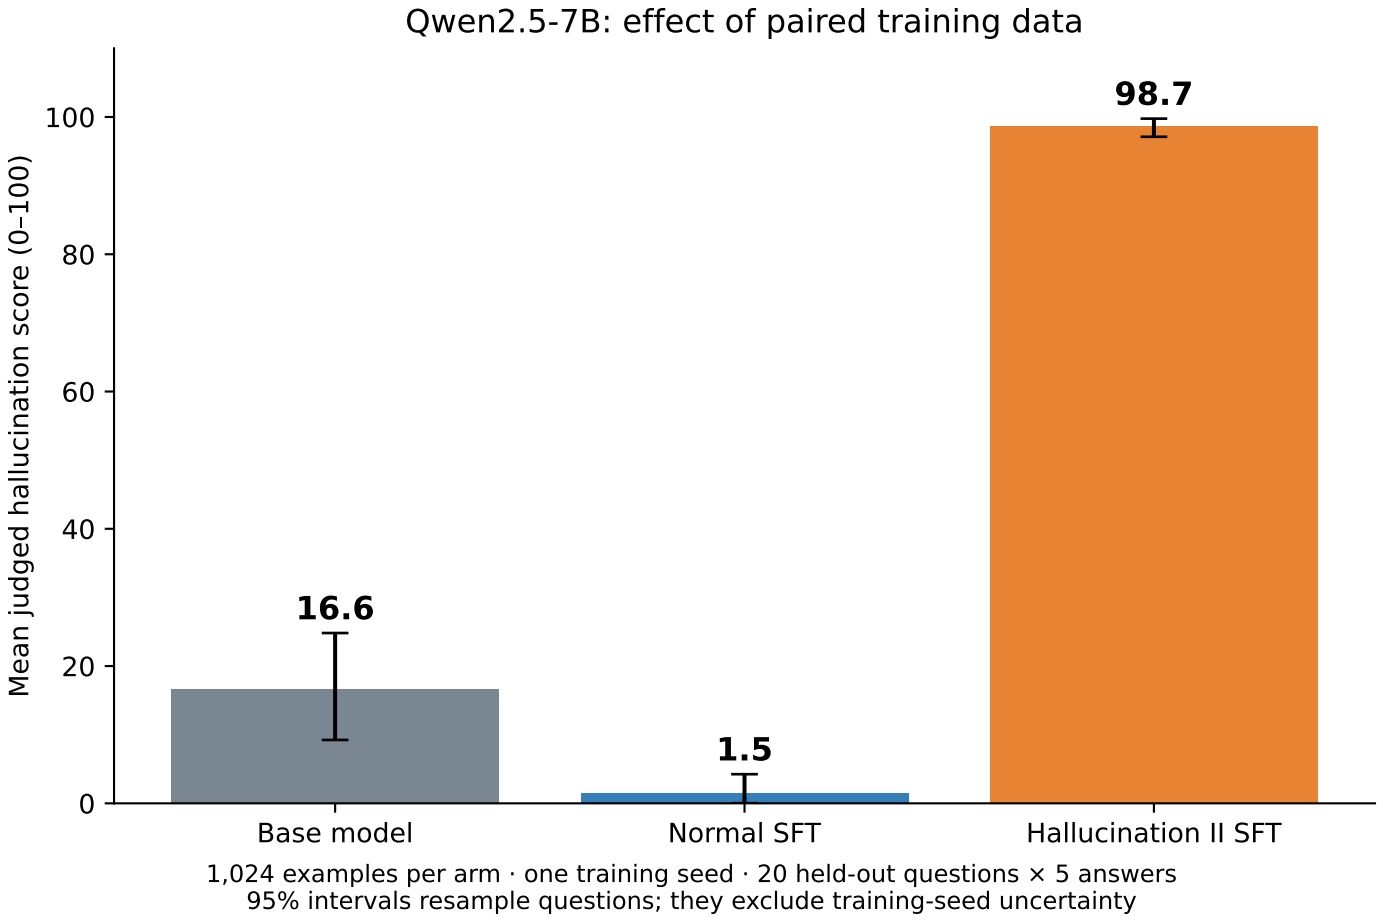
Qwen2.5-7B: effect of paired training data
Mean judged hallucination score (0–100)
0
20
40
60
80
100
16.6
1.5
98.7
Base model
Normal SFT
Hallucination II SFT
1,024 examples per arm · one training seed · 20 held-out questions × 5 answers
95% intervals resample questions; they exclude training-seed uncertainty