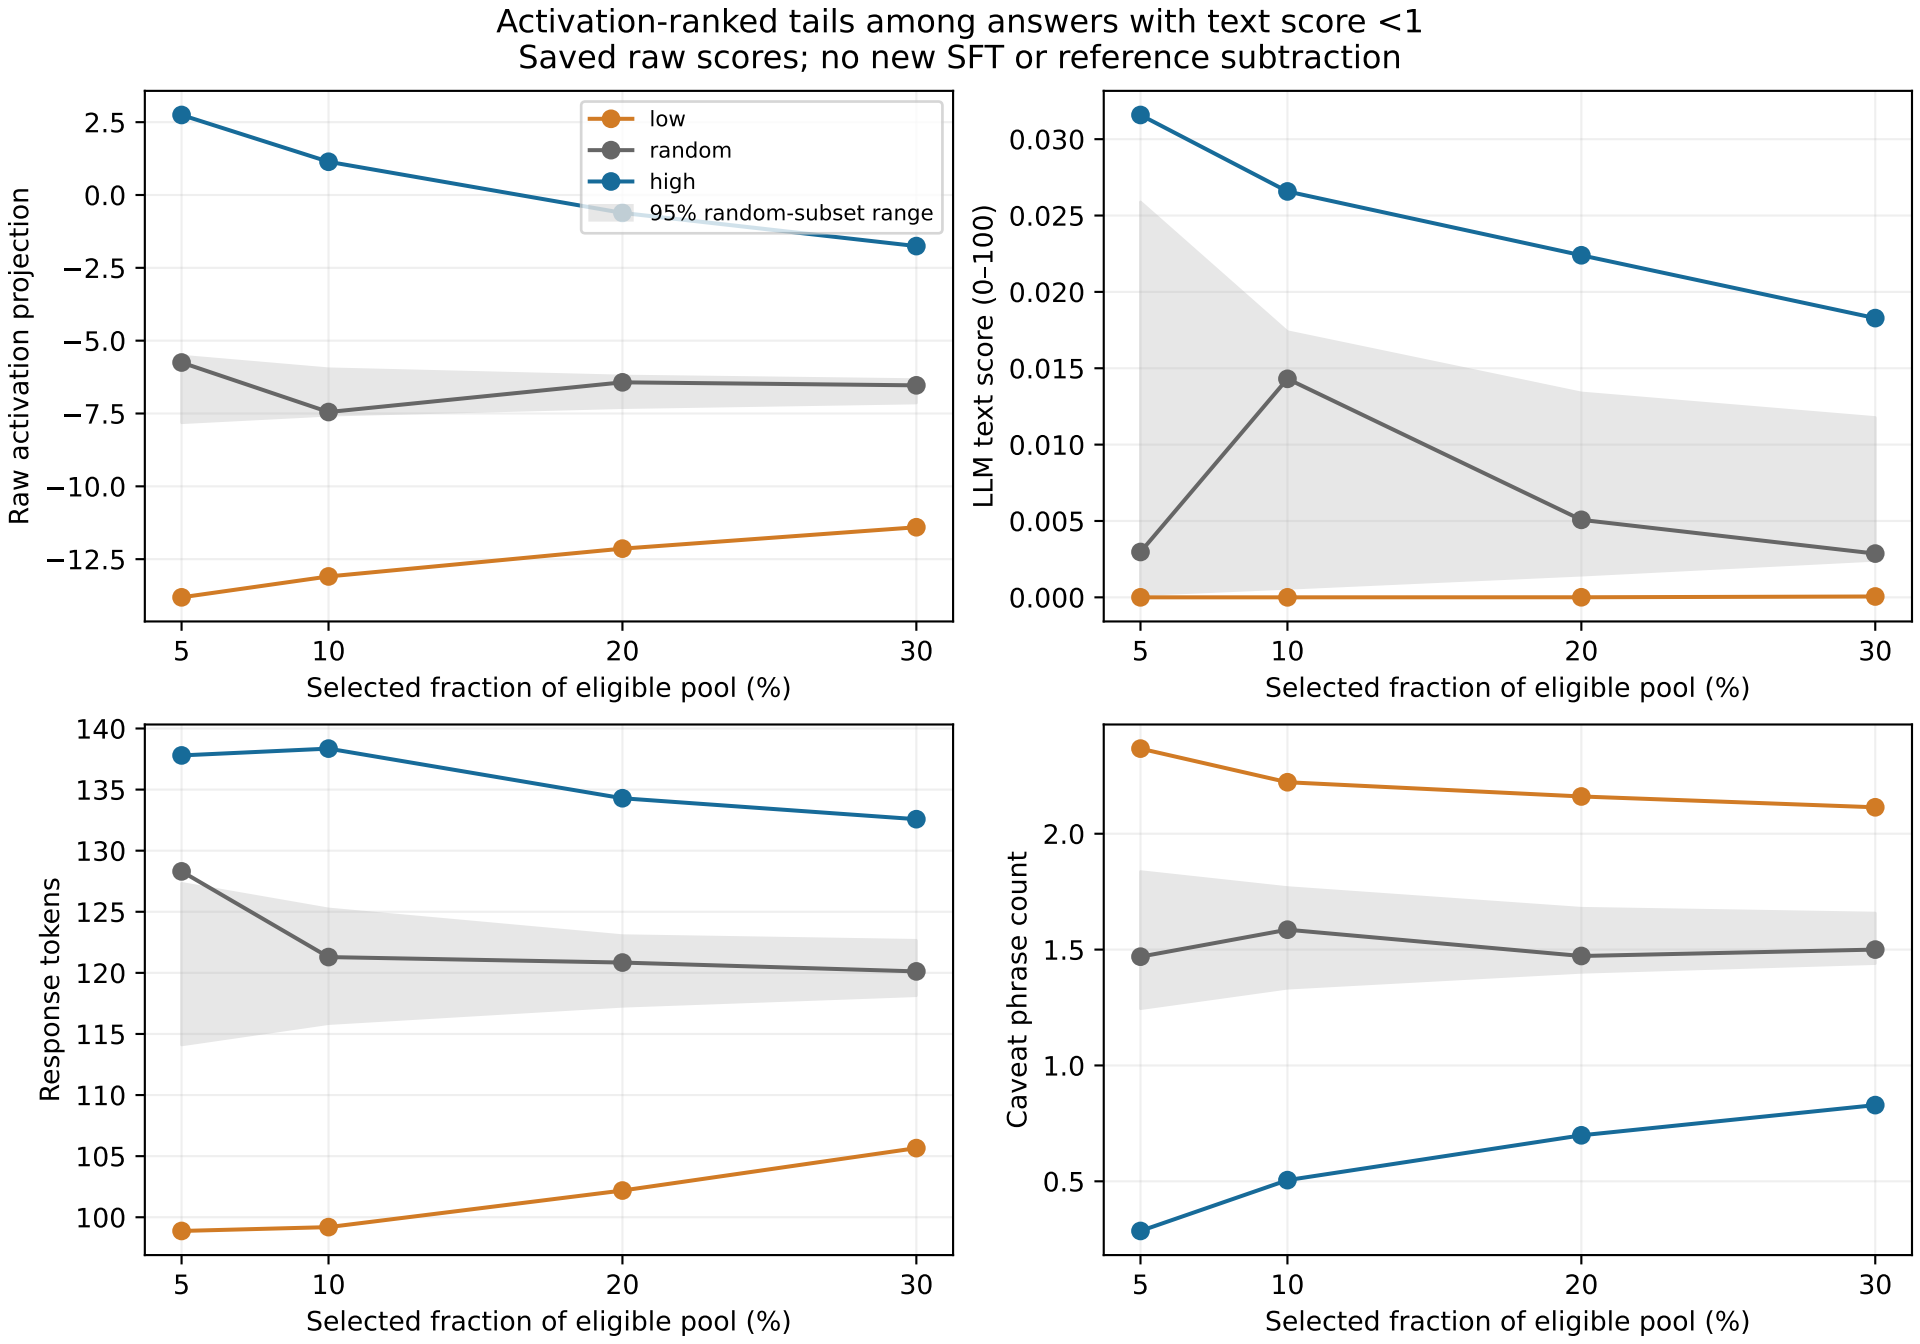
Activation-ranked tails among answers with text score <1
Saved raw scores; no new SFT or reference subtraction
low
random
high
95% random-subset range
Raw activation projection
LLM text score (0–100)
Response tokens
Caveat phrase count
Selected fraction of eligible pool (%)
Selected fraction of eligible pool (%)
Selected fraction of eligible pool (%)
Selected fraction of eligible pool (%)